

The Firtree Framework

A Dynamic Compilation System for Multi-Core Aware Image Processing

Dr Rich Wareham

Department of Engineering
University of Cambridge

Part I: Introduction

Outline

About Me

The Firtree Project

Requirements

Motivation

Existing Solutions

Plain C

GLSL

CoreImage

MapReduce

Lessons Learned

About Me

- ▶ Postdoc in the Signal Processing and Communications Laboratory¹, Department of Engineering, University of Cambridge
- ▶ Member of the Many Core Computing Group² in Cambridge
- ▶ Have a weakness for writing DSLs

¹<http://www-sigproc.eng.cam.ac.uk/>

²<http://www.many-core.group.cam.ac.uk/>

The Firtree Project

- ▶ Developed by one man (me!)
- ▶ Personal project developed after I wrote a Sobel filter for the 30th time
- ▶ Hosted at <http://firtree.org/>
- ▶ Aim is to allow me to write once, use many times
- ▶ An exercise in an intrinsically multi-core capable system

Requirements

- ▶ Implemented image processing algorithms in C
- ▶ Want to make use of multi-core hardware
- ▶ Don't want to be fussed with mechanism
- ▶ Moderate programming skills
- ▶ Have some knowledge of nodal image processing pipelines

Motivation

- ▶ Want to spend time developing *algorithms* not *methods*
- ▶ Multi-core computing is tricksy
 - ▶ Race conditions
 - ▶ Deadlocks
 - ▶ ... (many more)
- ▶ Don't want to worry about Image I/O
 - ▶ PNGs/JPEGs/TIFFs
 - ▶ Movie files
 - ▶ Camera sources
- ▶ Specify *intent* of algorithm
 - ▶ Abstract image sources/sinks
 - ▶ Independent processing steps
 - ▶ DAG-like layout

Outline

About Me

The Firtree Project

Requirements

Motivation

Existing Solutions

Plain C

GLSL

CoreImage

MapReduce

Lessons Learned

Plain C

- ▶ A *Lowest Common Denominator* language
- ▶ Fast (single threaded)
- ▶ Verbose

```
1 /* c-example.c */
2
3 #include <stdlib.h> /* malloc */
4 #include <...>      /* load() and save() */
5
6 int main(int argc, char* argv[])
7 {
8     unsigned char *input, *output;
9     int row, col;
10
11     input = load();
12     output = malloc(sizeof(unsigned char) * WIDTH * HEIGHT);
13     for(row = 0; row < 480; ++row)
14         for(col = 0; col < 640; ++col) {
15             float in =
16                 (float)(input[col + row*WIDTH]) / 256.f;
17             float out = sqrtf(in);
18             output[col + row*WIDTH] = (unsigned char)
19                 (out / 256.f);
20         }
21     save(output);
22     return 0;
23 }
```

Problems

- ▶ Complex and error prone
- ▶ Memory leaks easy (see example)
- ▶ Tends to lead to *ad hoc* solutions
 - ▶ Multiple channels?
 - ▶ Different pixel format?
- ▶ Requires external libraries
- ▶ Not easy to parallelise
 - ▶ Rely on auto vectorization per-CPU
 - ▶ Use extension like OpenMP for multi-CPU
- ▶ Abstractions lead to sub-optimal code

Outline

About Me

The Firtree Project

Requirements

Motivation

Existing Solutions

Plain C

GLSL

CoreImage

MapReduce

Lessons Learned

OpenGL Shader Language

- ▶ OpenGL 2.0 and above mandate programmable 'shader' support
- ▶ In a pixel shader, each pixel written to screen is generated via a small program from graphics-related inputs
 - ▶ Model co-ordinate
 - ▶ Texture co-ordinate
 - ▶ Surface normal
 - ▶ Light positions
 - ▶ Textures
 - ▶ etc...
- ▶ Geared to 3D realtime graphics

```
1 /* glsl-example.glsl */
2 uniform sampler2D inputTex;
3
4 vec4 shader_function()
5 {
6     vec2 pixel_coord = gl_MultiTexCoord[0];
7     vec4 input = texture2D(inputTex, pixel_coord);
8     return sqrt(input);
9 }
```

```
1 /* glsl-example.c */
2 setup_and_load_shader();
3 use_shader();
4 glBegin(GL_QUADS);
5     glTexCoord2f(0,0); glVertex2f(0,0);
6     glTexCoord2f(0,1); glVertex2f(0,1);
7     glTexCoord2f(1,1); glVertex2f(1,1);
8     glTexCoord2f(1,0); glVertex2f(1,0);
9 glEnd();
```

Problems

- ▶ Focussed on rendering, not image processing
- ▶ Lots of boilerplate code to set up rendering
- ▶ Tricky to render results anywhere other than screen
 - ▶ Requires extensions to earlier OpenGL versions
 - ▶ not all rendering formats supported

Outline

About Me

The Firtree Project

Requirements

Motivation

Existing Solutions

Plain C

GLSL

CoreImage

MapReduce

Lessons Learned

CoreImage

- ▶ Image processing framework from Apple³
- ▶ Transparently accelerated
- ▶ DAG-like rendering pipeline
- ▶ 'Minimal work' solution
 - ▶ Rendering only happens *Just In Time*
 - ▶ Only needed pixels are rendered

³<http://developer.apple.com/macosx/coreimage.html>

```
1 /* coreimage-builtins-example.m */
2
3 NSURL *url;
4 CIImage *source;
5 CIImage *result;
6 CIFilter *hueAdjust;
7
8 url = [NSURL fileURLWithPath:path];
9 source = [CIImage imageWithContentsOfURL:url];
10
11 hueAdjust = [CIFilter filterWithName:@"CIHueAdjust"];
12 [hueAdjust setValue:source forKey:@"inputImage"];
13 [hueAdjust setValue:[NSNumber numberWithFloat:2.094]
14                 forKey:@"inputAngle"];
15
16 result = [hueAdjust valueForKey:@"outputImage"];
```

```
1 /* coreimage-example.knl */
2
3 /* Apply a vertically varying haze removal filter. */
4 kernel vec4 myHazeRemovalKernel(sampler src,
5                                 __color color,
6                                 float distance,
7                                 float slope)
8 {
9     vec4 t;
10    float d;
11
12    d = destCoord().y * slope + distance;
13    t = unpremultiply(sample(src, samplerCoord(src)));
14    t = (t - d*color) / (1.0-d);
15
16    return premultiply(t);
17 }
```

Problems

- ▶ Only available on Mac OS X 10.5 and later
- ▶ A little too transparent
 - ▶ Cannot select backend explicitly
 - ▶ Only GPU and CPU backend implemented
- ▶ Only supports 1-to-1 pixel transforms

Outline

About Me

The Firtree Project

Requirements

Motivation

Existing Solutions

Plain C

GLSL

CoreImage

MapReduce

Lessons Learned

MapReduce

- ▶ Massively parallel database framework from Google
- ▶ Not graphics focussed *per se*
- ▶ Split processing into two sequential individually parallelisable processes
 - ▶ **Map:** Transform source data 1-to-1
 - ▶ **Reduce:** Compute overall statistic

Definition

- ▶ Define record set.
 - ▶ For example, each element has count for {male, total} children

$$\mathcal{R} = \{\{1, 2\}, \{2, 4\}, \{0, 1\}, \dots\}$$

- ▶ Wish to compute total number of female children
- ▶ Define *map* operation to calculate number of female children:

funct map(r) \equiv
 $(r_2 - r_1);$

- ▶ Define *reduce* operation to sum output from map

funct reduce(m, s) \equiv { s is pointer to a global state variable }
atomicinc(s, m);

Definition (cont...)

- ▶ MapReduce conceptually runs the following program:

```
begin
     $\mathcal{M} := \emptyset$ 
    for  $r \in \mathcal{R}$  do                                { loop 1 }
        appendtoset( $\mathcal{M}$ , map( $r$ ));
        od
     $s := 0$ 
    for  $m \in \mathcal{M}$  do                            { loop 2 }
        reduce(addressof( $s$ ),  $m$ );
        od
    end
```

Observations

- ▶ Supports efficient insertion of elements into set \mathcal{R}
 - ▶ Need only run one extra iteration of loops 1 and 2
- ▶ If inverse of *reduce* operation is defined removal of an element r from \mathcal{R} is also efficient
 - ▶ Need only call $\text{reduce}^{-1}(s, \text{map}(r))$
- ▶ Inherently parallel
- ▶ Not convenient for image processing operations
- ▶ Supports computing ensemble statistics
 - ▶ Not easy using accelerated frameworks such as GLSL and CoreImage

Lessons Learned

Suppose we wish to create the 'best of all worlds' system:

- ▶ Efficient API with minimal boilerplate
- ▶ Supports transparent acceleration where possible
 - ▶ But retain 'knobs' for advanced usage
- ▶ DAG-based pipeline workflow
- ▶ Image source/sink agnostic
- ▶ Inherently parallel
- ▶ Support 1-to-1 and ensemble operations

Part II: The Firtree API

Outline

The Firtree Image Model

API Bindings

GObject

Python

Samplers

Sampler Types

The Firtree Kernel Language

The Firtree Image Model

- ▶ Images are stored as pre-multiplied RGBA 4-tuples
 - ▶ Usually in the range (0,1) but not mandated to be so
 - ▶ Stored as single-precision floating point values
 - ▶ Optional double precision is an ultimate goal
 - ▶ Pre-multiplied so that common operations such as alpha over are efficient
- ▶ Images are conceptually infinite in extent
 - ▶ Do have a concept of a *non-transparent extent* however
- ▶ Pixel co-ordinates can be non-integer
- ▶ Integer co-ordinates refer to centre of the corresponding pixel
- ▶ Images consist of a *recipe* for how to calculate a pixel colour known as the *intent*

The Firtree API

- ▶ The Firtree API consists of a low-level GObject-based OO interface
 - ▶ The low-level API is easily bound to higher-level languages
 - ▶ Ships with Python API
 - ▶ Test-suite actually written in Python
- ▶ API is *intent* based
 - ▶ Record the intent of an image operation
 - ▶ DAG of input and processing nodes leading to a single output
 - ▶ No rendering is performed during specification of intent⁴
 - ▶ Intent compiled down to an intermediate representation
 - ▶ Pluggable backends render result on demand

⁴unless asked for!

Outline

The Firtree Image Model

API Bindings

GObject

Python

Samplers

Sampler Types

The Firtree Kernel Language

GObject

- ▶ Pure C object model API
- ▶ As nice as possible, given it is C
- ▶ Have to manage own memory
- ▶ Very low-level
- ▶ Used by a number of projects as their *glue layer*
- ▶ Designed to be bound to high-level languages

The Firtree GObject API

```
1 /* firtree-gobject.c */
2
3 #include <firtree/firtree.h>
4
5 /* ... */
6
7 FirtreeKernel *k = firtree_kernel_new();
8 const char[] src =
9     "kernel vec4 redKernel() {"
10    "    return vec4(1,0,0,1);"
11    "}";
12 firtree_kernel_compile_from_source(k, &src, 1, NULL);
13
14 g_object_unref(k);
```

Outline

The Firtree Image Model

API Bindings

GObject

Python

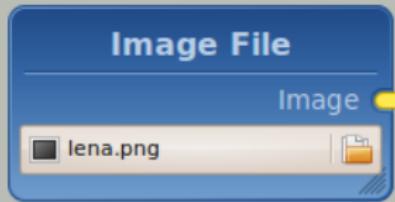
Samplers

Sampler Types

The Firtree Kernel Language

The Firtree Python API

- ▶ Automatically generated from GObject API
- ▶ Memory management taken care of by Python interpreter
- ▶ Language has object orientation syntactic sugar
- ▶ Greater SNR
- ▶ Interoperates with other Python libraries for image I/O, etc.
 - ▶ Don't re-invent the wheel
 - ▶ Let people use the tools they want



```
1 #!/usr/bin/env python
2
3 import cairo
4 import pyfirtree as ft
5
6 # Firstly, load the lena image
7 lena_surface = cairo.ImageSurface.create_from_png('lena.png')
8
9 # Create a sampler for the surface
10 lena_sampler = ft.CairoSurfaceSampler()
11 lena_sampler.set_cairo_surface(lena_surface)
12
13 # Create an output surface similar to the input
14 output_surface = cairo.ImageSurface(
15     cairo.FORMAT_ARGB32,
16     lena_surface.get_width(),
17     lena_surface.get_height() )
```

```
1 # Create a CPU render engine.  
2 engine = ft.CpuRenderer()  
3  
4 # Use the engine to write the input to the output.  
5 engine.set_sampler(lena_sampler)  
6 engine.render_into_cairo_surface(  
7     lena_sampler.get_extent(), # what area to render  
8     output_surface           # into what  
9 )  
10  
11 # Write the output  
12 output_surface.write_to_png('output.png')
```



Input



Output

Samplers

- ▶ Know how to generate the red, green, blue and alpha components of an output pixel given its location
- ▶ Has an associated transform which maps from image-space co-ordinates to sampler-space co-ordinates
 - ▶ Allows transformation of images via a *pass through* sampler
- ▶ Black boxes but may depend on other samplers
- ▶ Supports three primitive operations
 - ▶ Prepare: Called once before the render
 - ▶ Query: Called one or more times during the render
 - ▶ Finish: Called once after the render

Sampler Types

- ▶ A sampler is an abstract node
- ▶ Concrete implementations exist for various different images sources
- ▶ Buffer samplers sample from an in-memory image
 - ▶ Cairo surfaces, GdkPixbufs, images loaded by your own libraries (e.g. The PIL), etc.
 - ▶ Also a GStreamer plugin which uses a buffer sampler to provide video stream support
- ▶ Kernel samplers use a small piece of code to generate the output pixel
 - ▶ The code may, itself, sample from other samplers
- ▶ Advanced users can define their own samplers by implementing the three sampler primitives

The Firtree Kernel Language

- ▶ The most interesting samplers are *kernel samplers*
- ▶ A function which is called (conceptually) once per output pixel
- ▶ Returns the colour of a pixel given a *destination co-ordinate*
- ▶ Language is based on a subset of C
- ▶ Extended with vectorised types
 - ▶ $a + b \rightarrow$ element-wise addition
 - ▶ $a * b \rightarrow$ element-wise multiplication (not matrix multiplication)
- ▶ Very similar to GLSL or CoreImage kernels
 - ▶ Inspired by CoreImage
 - ▶ Originally implemented with the GLSL reference compiler

Part III: Examples

Outline

Desaturate

Sobel Filter

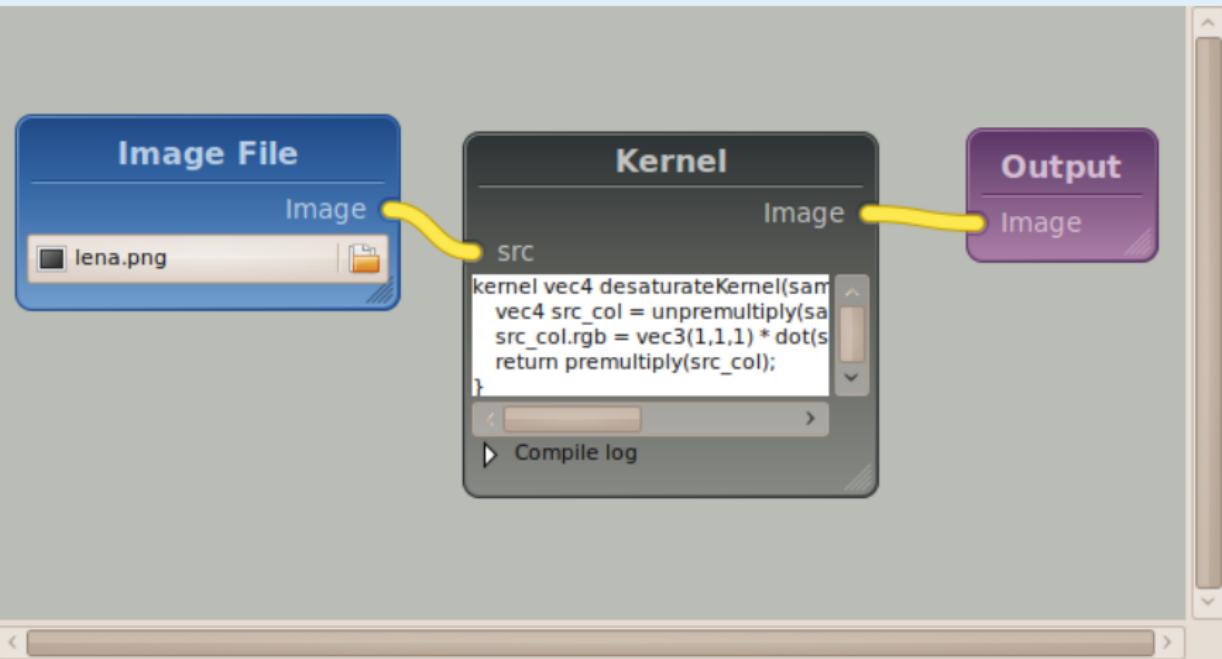
Mandelbrot Set

Reduce Kernels

Real-time Video Processing

Desaturate

- ▶ Convert an RGB image into the Y-component of a YUV image
- ▶
$$Y = 0.299R + 0.587G + 0.114B$$
- ▶ Use a simple kernel which uses the builtin **dot** function.



```
1 #!/usr/bin/env python
2
3 import cairo
4 import pyfirtree as ft
5
6 def compile_kernel(source):
7     """
8         A simple function which compiles a kernel, checks that
9         the compilation succeeded and returns a pair containing
10        the kernel and a sampler for it.
11    """
12
13    kernel = ft.Kernel()
14    kernel.compile_from_source(source)
15    if not kernel.get_compile_status():
16        print("Error compiling kernel:")
17        print("\n".join(kernel.get_compile_log()))
18        return
19
20    kernel_sampler = ft.KernelSampler()
21    kernel_sampler.set_kernel(kernel)
22
23    return (kernel, kernel_sampler)
```

```
1 # Firstly, load the lena image
2 lena_surface = cairo.ImageSurface.create_from_png('lena.png')
3
4 # Create a sampler for the surface
5 lena_sampler = ft.CairoSurfaceSampler()
6 lena_sampler.set_cairo_surface(lena_surface)
7
8 # Create a desaturate kernel.
9 (desat, desat_sampler) = compile_kernel("""
10     kernel vec4 desaturate(sampler src)
11     {
12         vec4 src_colour =
13             unpremultiply( sample(src, samplerCoord(src)) );
14         float luminance =
15             dot(src_colour, vec4(0.299,0.587,0.114,0));
16         return premultiply(
17             vec4(luminance,luminance,luminance,src.colour.a)
18         );
19     }
20 """")
21
22 # Wire the lena sampler into the desaturate kernel.
23 desat['src'] = lena_sampler
```

```
1 # Create an output surface similar to the input
2 output_surface = cairo.ImageSurface(
3     cairo.FORMAT_ARGB32,
4     lena_surface.get_width(),
5     lena_surface.get_height() )
6
7 # Create a CPU render engine.
8 engine = ft.CpuRenderer()
9
10 # Use the engine to render the output.
11 engine.set_sampler(desat_sampler)
12 engine.render_into_cairo_surface(
13     lena_sampler.get_extent(), # what area to render
14     output_surface           # into what
15 )
16
17 # Write the output
18 output_surface.write_to_png('output.png')
```



Input



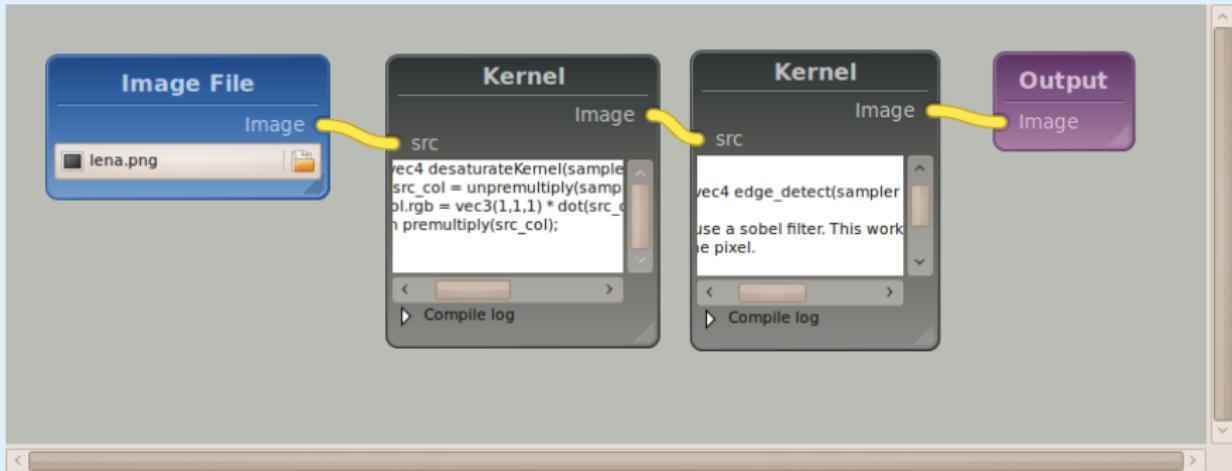
Output

Sobel Filter

- ▶ The Sobel Filter is a simple edge detector
- ▶ Convolves image with two kernels

$$I_x = I * \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} \quad \text{and} \quad I_y = I * \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- ▶ Computes edge response as $\sqrt{I_x^2 + I_y^2}$
- ▶ Need only append a kernel to the desaturate example above



```
1 # Create an edge detector kernel.  
2 (edge, edge_sampler) = compile_kernel("""  
3         // Return the (unpremultiplied) value of the pixel  
4         // at a particular offset.  
5  
6         ...  
7  
8         """)  
9  
1 edge['src'] = desat_sampler  
2  
3 # Use the engine to render the output.  
4 engine.set_sampler(edge_sampler)  
5 engine.render_into_cairo_surface(  
6     lena_sampler.get_extent(), # what area to render  
7     output_surface           # into what  
8 )  
9  
10 # Write the output  
11 output_surface.write_to_png('output.png')
```

```
1 // Return the (unpremultiplied) value of the pixel
2 // at a particular offset.
3 vec4 pixel_at(sampler src, vec2 offset)
4 {
5     return unpremultiply( sample( src,
6         samplerTransform( src, destCoord() + offset ) ) );
7 }
8
9 kernel vec4 edge_detect(sampler src)
10 {
11     // We use a sobel filter. This works on the 3x3 neighbourhood
12     // of the pixel.
13
14     // Get the neighbourhood.
15     vec4 p00 = pixel_at(src, vec2( -1, -1 ));
16     vec4 p10 = pixel_at(src, vec2( 0, -1 ));
17     vec4 p20 = pixel_at(src, vec2( 1, -1 ));
18     vec4 p01 = pixel_at(src, vec2( -1, 0 ));
19     vec4 p21 = pixel_at(src, vec2( 1, 0 ));
20     vec4 p02 = pixel_at(src, vec2( -1, 1 ));
21     vec4 p12 = pixel_at(src, vec2( 0, 1 ));
22     vec4 p22 = pixel_at(src, vec2( 1, 1 ));
```

```
1 float src_alpha = sample(src, samplerCoord(src)).a;
2
3 // Find the sobel response in X- and Y-directions.
4 vec4 dx = p00 - p20 + 2 * (p01 - p21) + p02 - p22;
5 vec4 dy = p00 - p02 + 2 * (p10 - p12) + p20 - p22;
6
7 // Calculate the magnitude of the response, we
8 // multiply by 0.25 so that the filter does not
9 // give a response > 1.
10 vec4 mag = 0.25 * sqrt( dx*dx + dy*dy );
11
12 return premultiply( vec4( mag.rgb, src_alpha ) );
13 }
```



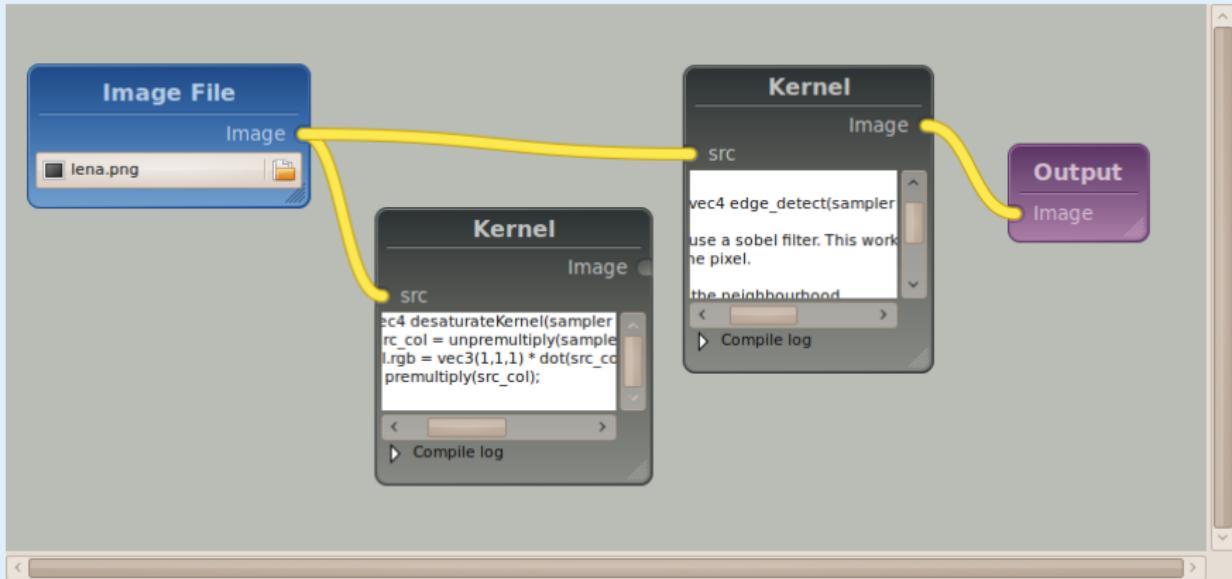
Input

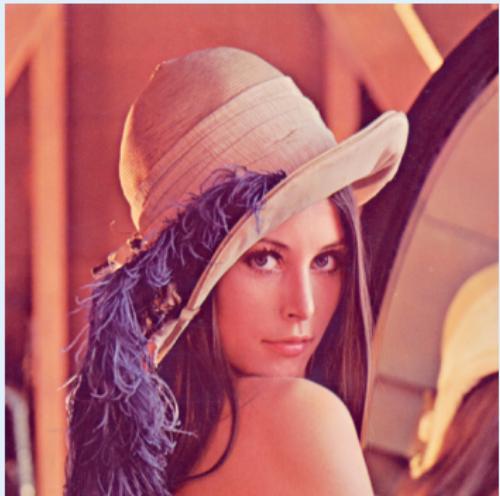


Output

Extending Sobel Filter

- ▶ The Sobel kernel we defined earlier used 4-way vector arithmetic
- ▶ We can extend the program to give a colour output easily by simply 'unplugging' the desaturate stage
- ▶ This requires precisely one line change:
1 `edge['src'] = lena_sampler`





Input



Output

```
1 # Create a kernel which computes the mandelbrot set
2 (mandelbrot, mandelbrot_sampler) = compile_kernel("""
3
4     vec2 complex_mul(vec2 a, vec2 b) {
5         return vec2(a.x*b.x - a.y*b.y, a.x*b.y + a.y*b.x);
6     }
7
7 kernel vec4 mandelbrot()
8 {
9     int max_iterations = 40;
10    vec2 c = destCoord(), z = c;
11
12    int num_its;
13    float mag_z_sq = dot(z,z);
14    for(num_its = 0;
15        (num_its < max_iterations) && (mag_z_sq < 4.0); ++num_its)
16        z = complex_mul(z,z) + c;
17        mag_z_sq = dot(z,z);
18
19    float output_val = num_its / max_iterations;
20    return vec4(output_val, output_val, output_val, 1);
21 }
```

```
1 output_surface = cairo.ImageSurface(  
2         cairo.FORMAT_ARGB32,  
3         1200, 800 )  
4 engine = ft.CpuRenderer()  
5 engine.set_sampler(mandelbrot_sampler)  
6 engine.render_into_cairo_surface(  
7         (-2, -1, 3, 2), # what area to render  
8         output_surface # into what  
9     )  
10 output_surface.write_to_png('mandelbrot.png')
```

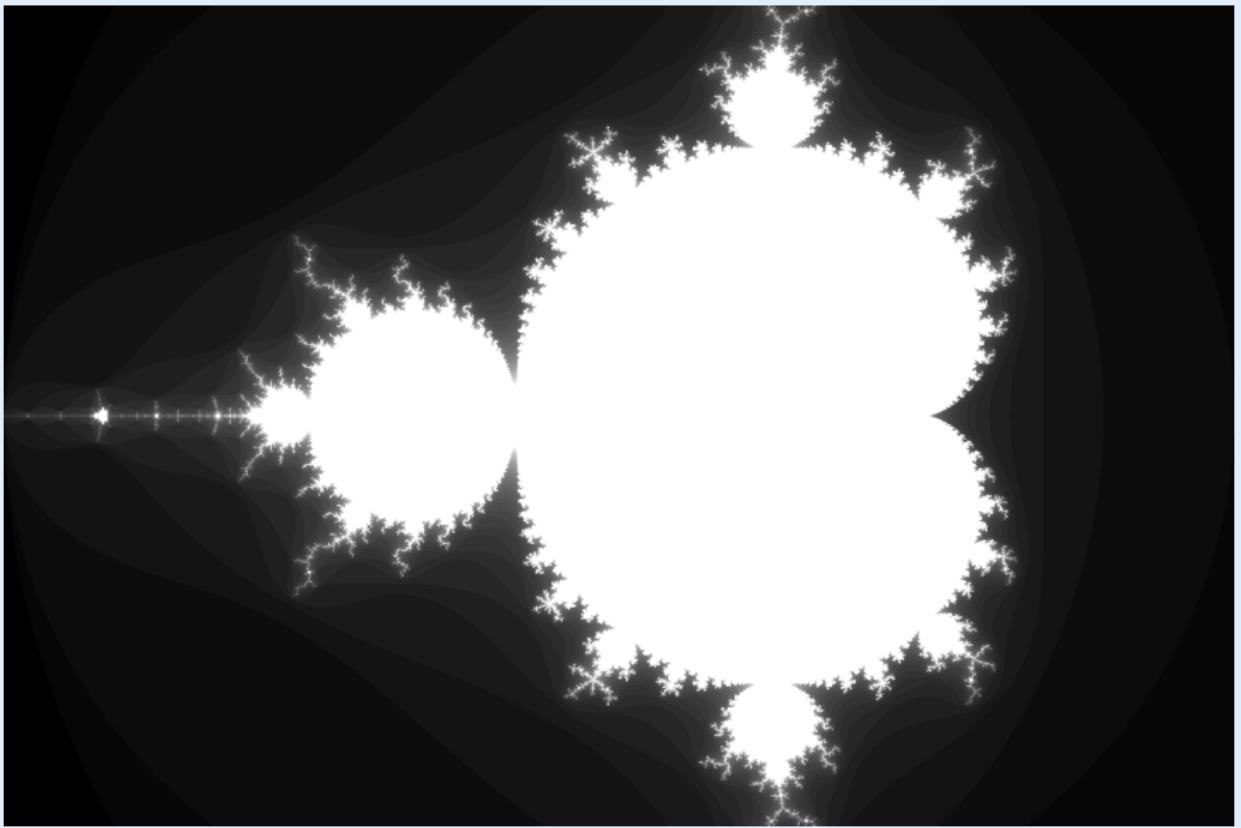


Figure: Output of `mandelbrot.py`

Reduce Kernels

- ▶ Like normal image kernels except that they have no output
 - ▶ Kernel function returns **void**
 - ▶ They cannot be associated with a sampler
- ▶ Are conceptually called once per-pixel but in undefined order
- ▶ May make use of one extra builtin: **emit()**
 - ▶ Appends a **vec4** value to a global list
- ▶ Somewhat misnamed; it is more of a selective filter + map

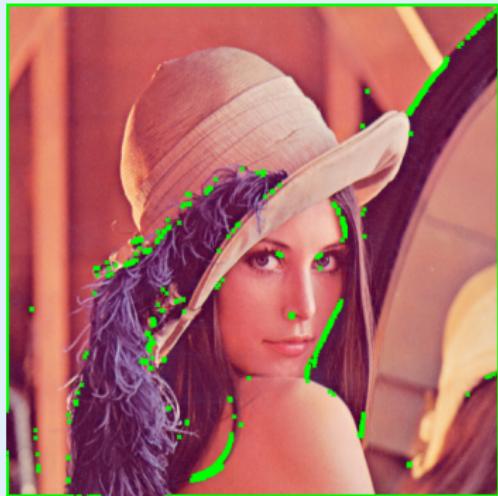
```
1 def compile_reduce_kernel(source):
2     kernel = ft.Kernel()
3     kernel.compile_from_source(source)
4     if not kernel.get_compile_status():
5         print("Error compiling kernel:")
6         print("\n".join(kernel.get_compile_log()))
7         return
8     return kernel
```

```
1 # Create a reduce kernel which looks for areas of high red.
2 reduce = compile_reduce_kernel("""
3     kernel __reduce void high_red(sampler src)
4     {
5         vec4 src_colour =
6             unpremultiply( sample(src, samplerCoord(src)) );
7         if( src_colour.r - max(src_colour.g, src_colour.b) > 0.1 )
8         {
9             emit( vec4( destCoord(), 0, 0 ) );
10        }
11    }
12 """)
13
14 reduce['src'] = edge_sampler
15
16 # Create a CPU reduce engine.
17 engine = ft.CpuReduceEngine()
18 engine.set_kernel(reduce)
```

```
1 # Find the matching points
2 edge_points = engine.run(
3     lena_sampler.get_extent(),           # what area of the input
4     lena_surface.get_width(),           # how many horizontal samples
5     lena_surface.get_height()           # how many vertical samples
6 )
7 print('%i edge points found' % len(edge_points))
8
9 # Draw them on the surface
10 cr = cairo.Context(lena_surface)
11 cr.set_source_rgba(0,1,0,1)
12 for point in edge_points:
13     cr.move_to(point[0] - 2, point[1] - 2)
14     cr.line_to(point[0] + 2, point[1] + 2)
15     cr.move_to(point[0] + 2, point[1] - 2)
16     cr.line_to(point[0] - 2, point[1] + 2)
17     cr.stroke()
18
19 # Write the output
20 lena_surface.write_to_png('output.png')
```

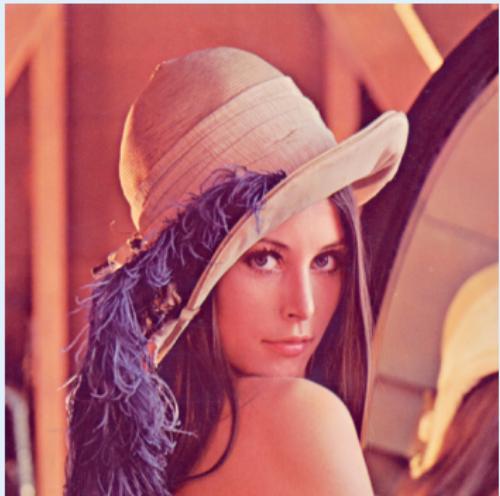


Input

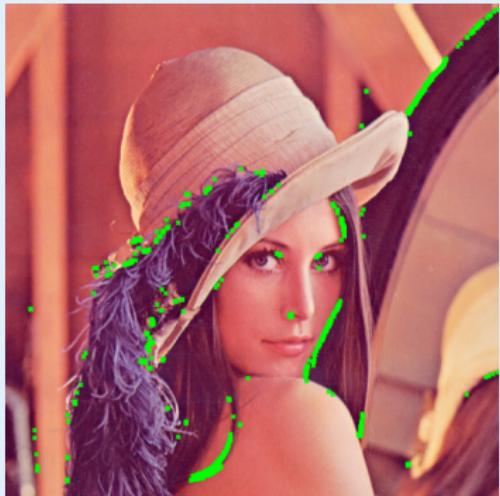


Output

```
1 valid_extents = lena_sampler.get_extent()
2 valid_extents = (
3     valid_extents[0] + 1, valid_extents[1] + 1,
4     valid_extents[2] - 2, valid_extents[3] - 2 )
5
6 # Find the matching points
7 edge_points = engine.run(
8     valid_extents,           # what area of the input to ren-
9     int(valid_extents[2]),   # how many horizontal samples
10    int(valid_extents[3]),   # how many vertical samples
11    )
12
13 # Draw them on the surface
14 cr = cairo.Context(lena_surface)
15 cr.set_source_rgba(0,1,0,1)
16 for point in edge_points:
17     cr.move_to(point[0] - 2, point[1] - 2)
18     cr.line_to(point[0] + 2, point[1] + 2)
19     cr.move_to(point[0] + 2, point[1] - 2)
20     cr.line_to(point[0] - 2, point[1] + 2)
21     cr.stroke()
22
23 # Write the output
24 lena_surface.write_to_png('output.png')
```



Input



Output

Real-time Video Processing

- ▶ Sampler is sufficiently general that we can make a video source a sampler
- ▶ Have integrated Firtree with the GStreamer project
 - ▶ GStreamer allows one to create 'pluggable' modules that can decode video, camera streams, etc
- ▶ Create a GStreamer *sink* node which appears as a Firtree sampler

Kernel

```

vec2 complex_mult(vec2 a, vec2 b) {
    return vec2(a.x*b.x - a.y*b.y, a.x*b.y + a.y*b.x);
}

kernel vec4 mandelbrot() {
    int max_iterations = 40;
    vec2 c = (destCoord() - vec2(800,400)) / 200;
    vec2 z = c;

    int num_its;
    float mag_z_sq = dot(z,z);
    for(num_its = 0; (num_its < max_iterations) && (mag_z_sq < 4.0); ++num_its) {
        z = complex_mult(z,z) + c;
        mag_z_sq = dot(z,z);
    }

    float output_val = num_its / max_iterations;
    return premultiply(vec4(1,output_val,1,min(1.10*output_val)));
}

```

Kernel

Image

src

```

kernel vec4 falseColor(sampler src) {
    vec4 src_col =
        unpremultiply(samplerColor(src, samplerCoord(src)));
    float src_color_a = src_col.a;
    vec3 color = 0.5 - 0.5 * cos(vec3(
        (2 * src_color_a * 3.14159),
        (4 * src_color_a * 3.14159),
        (8 * src_color_a * 3.14159)));
    return premultiply(vec4(color, src_col.a));
}

```

Compile log

Kernel

Image

top
bottom

```

kernel vec4 overKemel(sampler top, sampler bott)
vec4 top_col = sample(top, samplerCoord(top));
vec4 bottom_col = sample(bott, samplerCoord(bott));
return (1-top_col.a)*bottom_col + top_col;
}

```

Compile log

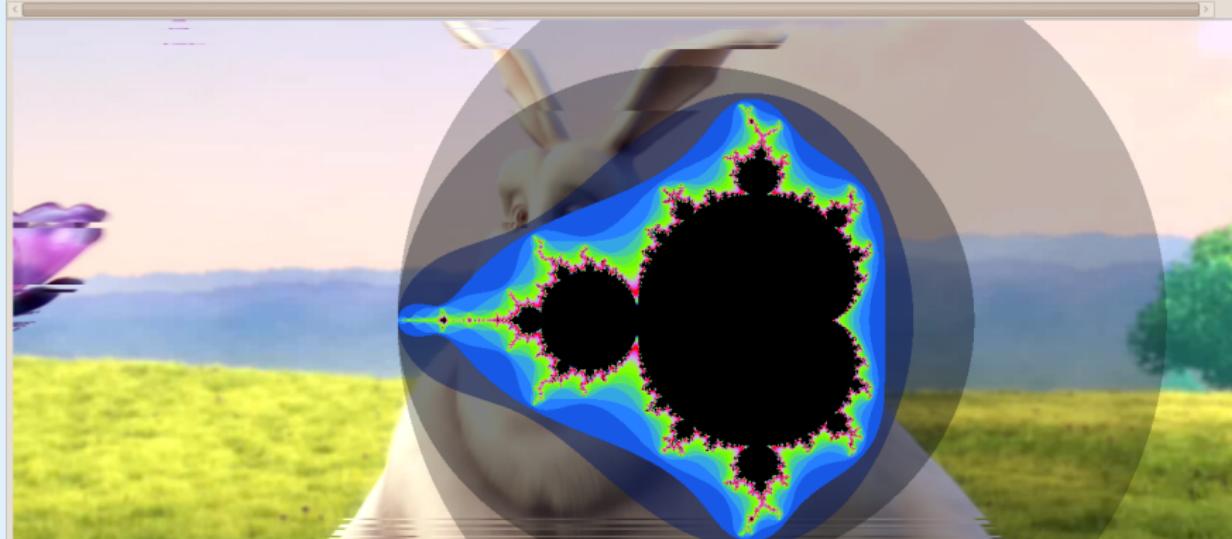
Output

Image

Video File

Video

big_buck_bunny_720p...



Part IV: Implementation

Outline

Compilation Process

LLVM

Sampler Linking

Backends

Multi-core CPU

The future: GPU

Compilation Process

- ▶ Frontend compilation
 - ▶ Each sampler can represent its sampling action via a common Intermediate Representation (IR)
 - ▶ IR has a number of low-level intrinsic functions
 - ▶ Sampling from image buffers, arithmetic operations, transcendental/trigonometric functions, etc.
- ▶ Backend execution
 - ▶ On first render, backend JIT compiles IR into native code
 - ▶ Backend distributed work units to each compute unit
 - ▶ e.g. CPU backend uses one thread per CPU which render blocks of output independently

Compilation Process (cont...)

- ▶ Each sampler generates IR for its action
- ▶ The IR linker automatically combines each sampler with the IR for its dependent samplers to produce a single block of IR
- ▶ The linking stage and those that precede it are backend independent
- ▶ The same IR may be passed to different backends if necessary

LLVM

- ▶ Make use of the IR from the LLVM project⁵
- ▶ Provides an in-memory and human-readable variant
- ▶ Includes a suite of modern optimisation stages
- ▶ Includes native vector-aware backends for most modern CPUs
- ▶ nVidia have a backend which targets their compute-capable graphics cards

⁵<http://llvm.org/>

LLVM Example

- ▶ Use the simple identity pipeline example

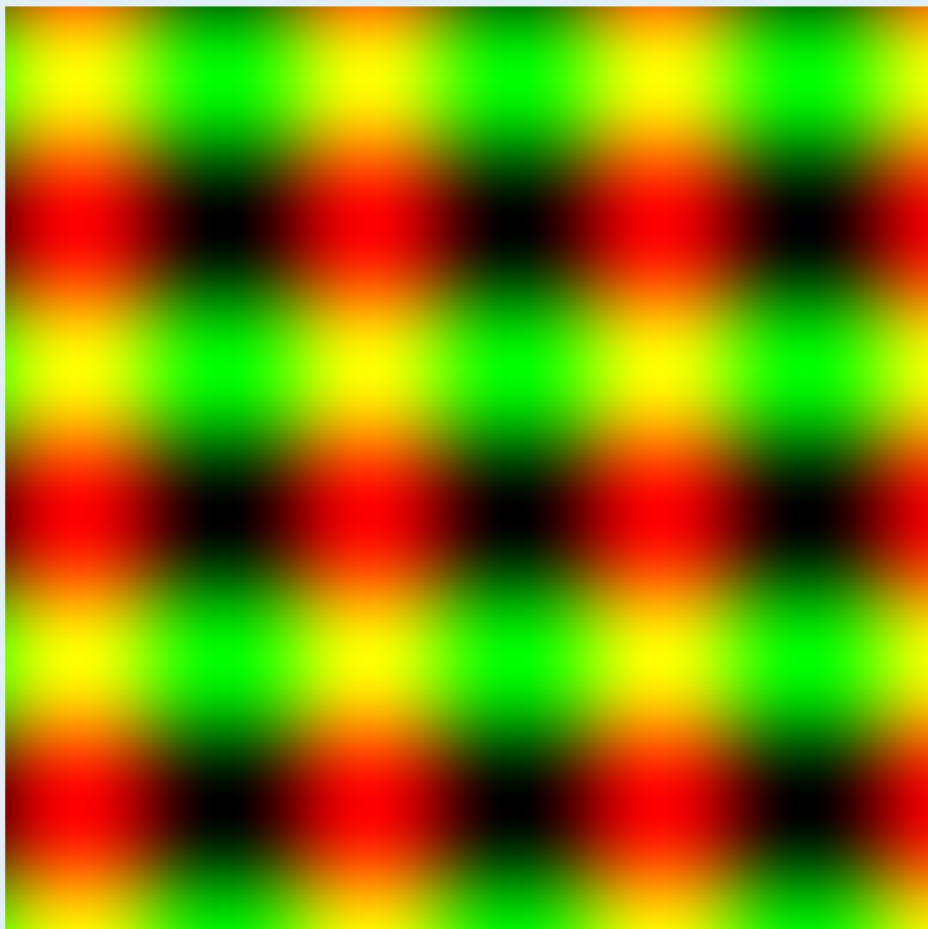
```
1 lena_sampler = ft.CairoSurfaceSampler()  
2 lena_sampler.set_cairo_surface(lena_surface)
```

- ▶ Sampler generates following IR

```
1 ; Builtin IR intrinsic  
2 declare <4 x float> @sample_image_buffer_nn(  
3     i8*, i32, i32, i32, <4 x float>)  
4  
5 ; Sampler function itself  
6 define <4 x float>  
7     @sampler_dc2e6e61_1920_08f4_898f_2456c40f1e6b(<4 x float>) {  
8 entry:  
9     %rv = call <4 x float> @sample_image_buffer_nn(  
10        i8* inttoptr (i64 -1270542328 to i8*),  
11        i32 13, i32 512, i32 512,  
12        i32 2048, <4 x float> %0) ; <<4 x float>> [#uses=1]  
13     ret <4 x float> %rv  
14 }
```

```
1 kernel vec4 sinKernelExample()
2 {
3     return vec4( 0.5 + 0.5 * sin( destCoord() ), 0, 1 );
4 }

1 ; IR builtin
2 declare <4 x float> @sin_v2(<4 x float>) nounwind readnone
3
4 ; Sampler function proper
5 define <4 x float>
6     @sampler_35014ecb_ea33_8994_f1ed_88ceb0e101b7(<4 x float>) {
7 entry:
8     %tmp.i = call <4 x float> @sin_v2(<4 x float> %0) nounwind
9     %tmp3.i = fmul <4 x float> %tmp.i, <float 5.000000e-01,
10      float 5.000000e-01, float 0.000000e+00, float 0.000000e+00>
11     %tmp6.i = fadd <4 x float> %tmp3.i, <float 5.000000e-01,
12      float 5.000000e-01, float 0.000000e+00, float 0.000000e+00>
13     %tmp13.i = insertelement <4 x float> %tmp6.i,
14      float 0.000000e+00, i32 2
15     %tmp14.i = insertelement <4 x float> %tmp13.i,
16      float 1.000000e+00, i32 3
17     ret <4 x float> %tmp14.i
18 }
```



Sampler Linking

- ▶ Samplers may call other samplers

```
1 kernel vec4 sinexample(sampler src)
2 {
3     vec2 dc = destCoord() + 5 * sin(destCoord() / 10);
4     return sample(src, samplerTransform(src, dc));
5 }
```

- ▶ Kernel is compiled into IR which makes use of sampler functions

```
1 kernel vec4 sinKernelExample()
2 {
3     return vec4( 0.5 + 0.5 * sin( destCoord() ), 0, 1 );
4 }

1 ; Intrinsics
2 declare <4 x float> @sin_v2(<4 x float>) nounwind readnone
3 declare <4 x float> @samplerTransform_sv2(i32, <4 x float>) nounwind readnone
4 declare <4 x float> @sample_sv2(i32, <4 x float>) nounwind readnone
5
6 ; Kernel function
7 define <4 x float>
8     @kernel_4bcd26e6_ec72_7454_f967_7b7d92125bba(
9         <4 x float> %__this__destCoord, i32 %src) nounwind readnone {
10    entry:
11        %tmp3 = fdiv <4 x float> %__this__destCoord, <float 1.000000e+01,
12            float 1.000000e+01, float 0.000000e+00, float 0.000000e+00>
13        %tmp4 = call <4 x float> @sin_v2(<4 x float> %tmp3)
14        %tmp8 = fmul <4 x float> %tmp4, <float 5.000000e+00, float 5.000000e+00,
15            float 0.000000e+00, float 0.000000e+00>
16        %tmp9 = fadd <4 x float> %tmp8, %__this__destCoord
17        %tmp13 = call <4 x float> @samplerTransform_sv2(i32 %src, <4 x float> %tmp9)
18        %tmp14 = call <4 x float> @sample_sv2(i32 %src, <4 x float> %tmp13)
19        ret <4 x float> %tmp14
20 }
```

Linker Output

Setting lena as `src` causes the linker to inline calls to the sampler:

```
1 ; Builtins
2 declare <4 x float> @sin_v2(<4 x float>) nounwind readnone
3 declare <4 x float> @sample_image_buffer_nn(i8*, i32, i32, i32, <4 x float>)
4
5 ; Sampler function
6 define <4 x float>
7 @sampler_8e662b89_bf6d_26f4_412b_61c73e556e92(<4 x float>) {
8 entry:
9 %tmp3.i = fdiv <4 x float> %0, <float 1.000000e+01, float 1.000000e+01,
10    float 0.000000e+00, float 0.000000e+00>
11 %tmp4.i = call <4 x float> @sin_v2(<4 x float> %tmp3.i) nounwind
12 %tmp8.i = fmul <4 x float> %tmp4.i, <float 5.000000e+00, float 5.000000e+00,
13    float 0.000000e+00, float 0.000000e+00>
14 %tmp9.i = fadd <4 x float> %tmp8.i, %0
15 %rv.i.i.i = call <4 x float> @sample_image_buffer_nn(
16    i8* inttoptr (i64 -1269264376 to i8*), i32 13, i32 512, i32 512, i32 2048,
17    <4 x float> %tmp9.i) nounwind
18 ret <4 x float> %rv.i.i.i
19 }
```



Backends

Backends are responsible for platform specific tasks:

- ▶ Compiling LLVM IR to native code
- ▶ Implementing the intrinsics
- ▶ Distributing rendering onto multiple cores

Outline

Compilation Process

LLVM

Sampler Linking

Backends

Multi-core CPU

The future: GPU

The CPU backend

- ▶ Wraps the sampler IR in LLVM IR which implements an inner rendering loop
 - ▶ Similar to the plain C example at the start of the talk
- ▶ Maintains one worker thread per-CPU
- ▶ Divides rendering task into small chunks
- ▶ Each worker thread gets a new chunk from a global queue
- ▶ Easier parts of the image do not cause threads to become idle

The generated CPU assembly (x86 in this case) can be quite daunting

Outline

Compilation Process

LLVM

Sampler Linking

Backends

Multi-core CPU

The future: GPU

The future: GPU

- ▶ Post Firree re-write GPU renderer has not been re-implemented
- ▶ A proof-of-concept LLVM IR to OpenCL/CUDA converter exists
- ▶ Not yet fully implemented
 - ▶ Very small number of intrinsics
 - ▶ Does not execute the code it compiles
- ▶ Waiting for an excuse to implement it
- ▶ Some interesting issues arise
 - ▶ Backend will also have to maintain buffer cache on compute card
 - ▶ Want to create some render *queue* so that entire rendering pipeline becomes asynchronous

The past: GPU

- ▶ An earlier version of Firtree had a GPU backend
- ▶ Was based on OpenGL + GLSL
- ▶ The 'unique' way OpenGL contexts are managed caused it to be brittle
- ▶ Was very fast indeed
- ▶ Even compute-intensive kernels required negligible CPU
- ▶ Some pretty videos were generated

Questions?